Annotated Explanatory Edition

A Spectral–Analytic Framework for P vs. NP

Brendan Philip Lynch

November 2025

Abstract

This annotated edition explains, contextualizes, and interprets the ideas of the *Spectral–Analytic Separation of P and NP* framework. It follows the explanatory format of the earlier Riemann–Hypothesis edition: each section of the main paper is accompanied by commentary for three audiences (technical, interdisciplinary, and public). The document does not assert or verify any proof of $P \neq NP$; it clarifies definitions, motivations, and the analytical–numerical structure of the model so that readers from mathematics, physics, and computer science can follow the reasoning and reproduction protocol.

Symbol Table — plain and technical explanations

Symbol / Term	Plain English	Technical Description
Φ_n	Circuit encoder: turns a Boolean circuit into numerical data	Polynomial-time mapping from circuits to Jacobi coefficients (a_k, b_k) of a discrete Schrödinger operator.
J	The "Jacobi matrix" representation of a circuit	Tri-diagonal real symmetric matrix whose entries define the discrete spectral problem.
$V(x)$	Potential function	Continuous half-line potential reconstructed via the Gelfand–Levitan–Marchenko (GLM) equation.
GLM transform	Reverse-engineering tool	Integral transform recovering $V(x)$ from spectral data; a continuous analogue of inverting J.
NCH	No-Compression Hypothesis	Information-theoretic postulate that exponential witness sets of NP problems cannot be encoded into polynomially many bounded-precision reals.
L^1-integrable potential	"Finite-strength" potential	A potential $V(x)$ with $\int_0^\infty \lvert V(x) \rvert dx < \infty$; ensures well-posed self-adjoint dynamics.

1 Introduction and Motivation

The P versus NP problem is a central challenge in computer science. This work transfers the problem from discrete complexity theory to continuous analysis via inverse spectral theory, leveraging the Gelfand–Levitan–Marchenko (GLM) transform. The core idea is to encode the complexity of a circuit C (specifically, its number of accepting witnesses) into the spectral measure of a Jacobi matrix J = n(C), which, in turn, maps to the L1-integrability of a continuous potential V (x) on the half-line. The crucial conditional step is the No-Compression Hypothesis (NCH), which dictates that the required information content of an NP-complete problem cannot be stored in an 1-summable Jacobi sequence with polynomial precision and length.

Definition 1.1 (Bit model / complexity conventions). All time bounds poly(n) in this paper are measured in the standard multi-tape Turing machine bit-complexity model: inputs are encoded in binary; arithmetic on integers of O(b(n)) bits has cost ˜O(b(n)) per operation using standard multiplication algorithms; rational output of b(n) bits counts toward the output representation cost. If instead a word-RAM model is preferred, add an explicit conversion clause; the present statements use the Turing-bit model.

(Review Committee / Technical — Blue)

This section outlines the spectral mapping from circuits to potentials and introduces the NCH as the key conditional. Emphasize the transfer via GLM and the 1 L1 correspondence (Lemma 3.1). Be prepared to detail how E1–E4 ensure injectivity and computability, and note the bit-model precision to avoid word-RAM ambiguities.

(Interdisciplinary Professionals — Green)

Think of this as bridging computer science and physics: circuit complexity becomes a quantum-like potential landscape. P problems have "finite energy" (L1-integrable), while NP require infinite, mirroring chaotic vs. ordered systems. The NCH is like saying you can't compress exponential data into polynomial space without loss.

(General Public / Press — Orange)

Imagine turning a tough puzzle (like NP) into a hilly path for a ball to roll on. Easy puzzles (P) have finite hills, but hard ones go on forever. The "no-compression" idea means you can't stuff all the puzzle pieces into a small bag—they spill out, proving the puzzles are fundamentally different.

2 Formal Conditional Theorem

We first state the core conditional theorem that formalizes the separation.

2.1 Theorem (Conditional Analytic Separation)

Theorem 2.1. Let n be a family of computable circuit-to-Jacobi encodings satisfying properties (E1)–(E4) below. Suppose the following hold:

1. No-Compression Hypothesis (NCH)/Packing Lower Bound (PLB). There exist constants C, ¿ 0 and functions b(n) = poly(n), T(n) = poly(n), such that the injective poly(n)-bit encoding of 2n NP-complete witnesses requires a Jacobi sequence length m(n) n (super-polynomial), satisfying the packing inequality:

m(n) k=1 log2 (1 + C2b(n)

T(n)k2

) n.

2. GLM Stability and Recovery (GSR). If a Jacobi matrix J has entries with b(n)-bit rational precision and k=1(—ak1—+—bk—) ¡ (i.e., V L1), then the GLM inverse problem reconstructs J (and hence the discrete signature SC) in poly(n) time and poly(n) bit complexity to sufficient precision.

Then, under (A) and (B), for every NP-complete language L, the continuous potentials VNP corresponding to accepting-circuits for L are not L1-integrable, while those for P are L1-integrable; hence P = NP.

(Review Committee / Technical — Blue)

The theorem hinges on NCH implying super-polynomial m(n), leading to 1 divergence for NP. Highlight GSR's role in recovery complexity and prepare asymptotic expansions for the packing sum. Note E4's necessity to prevent "cheating" via faster decay.

(Interdisciplinary Professionals — Green)

This is the main result: under NCH, NP "overflows" the L1 bound, like a system with unbounded energy in physics. E1–E4 ensure the mapping is fair, similar to constraints in signal processing where decay rates limit information density.

(General Public / Press — Orange)

The big claim: if you can't cram hard puzzle info into a small math box (NCH), then hard puzzles need infinite landscapes, unlike easy ones. The rules (E1–E4) make sure the translation is honest—no shortcuts.

3 Explicit Encoding Properties (E1)–(E4)

The encoding n must satisfy the following explicit constraints:

1. Computability & Canonical Form. n is computable in poly(n) time and maps circuits to rational Jacobi entries with known rational denominators, bounded by 2b(n), where b(n) = poly(n) (e.g., b(n) = n2).

2. Local-Amplitude Bound. The per-index amplitude window Ik for —ak 1— and —bk— is Ik [0, C/(T(n)k2)]. This enforces the 1 summability constraint on the differences from the identity matrix for V L1.

3. Recovery Uniqueness (Injectivity). For any two circuits C = C, the resulting Jacobi matrices n(C) and n(C) differ by a minimum index-wise separation k such that k ¿ 21b(n) for k m(n). This proves injectivity under poly(n) bit rounding.

4. Index-Role Invariance (Decay Constraint). The encoding n is restricted to obey the canonical decay rate of O(1/k2). Any encoding that attempts to concentrate all 2n bits into O(n) early coordinates (violating 1-summability) must break the poly(n) precision bound (E1) to satisfy the n packing condition.

(Review Committee / Technical — Blue)

Each property ensures the encoding Φ_n behaves mathematically well: computability, bounded precision, injectivity, and a canonical $1/k^2$ decay. They prevent pathological encodings that would break the spectral machinery.

(Interdisciplinary Professionals — Green)

These are the "engineering tolerances" that keep the translation from circuits to spectra stable.

(General Public / Press — Orange)

Think of them as quality-control rules: numbers must fade smoothly so the system stays well behaved.

4 Detailed Proof Skeleton and Lemmas

3.1 Lemma 3.1: Discrete 1 Continuous L1 Transfer

Based on results in inverse spectral theory (see e.g., work of Gesztesy and Simon), the Jacobi matrix J = ak, bkk1 corresponds to a continuous half-line potential V (x) in L1([0,)) if and only if the coefficients satisfy the discrete 1 condition: k=1(—ak1—+—bk—) ¡ M ¡ . Conversely, if the discrete sum diverges, V (x) is not L1-integrable.

1. P Case: For P problems, n bits of information are sufficient to encode the complexity, requiring a length m(n) = O(log n). Under the O(1/k2) decay, the 1 sum is bounded: MP = O(1), hence VP L1 = O(1).

2. NP Case (Under NCH): The NCH (PLB) forces m(n) to be super-polynomial. Since the decay is fixed at O(1/k2), the super-polynomial length forces the 1 norm to diverge: k=1(—ak 1— + —bk—) →. Hence, VNP L1 →.

(Review Committee / Technical — Blue)

This lemma connects discrete ℓ^1 summability of Jacobi coefficients to continuous L^1 integrability of the potential.

(Interdisciplinary Professionals — Green)

Finite total deviation in the discrete model means finite total energy in the continuous one.

(General Public / Press — Orange)

If you can add up all the bumps on the line and get a finite number, the system is tame; otherwise it's wild.

5 GLM Reconstruction Complexity and Stability

The GLM reconstruction is performed via Nyström discretization (Assumption B).

• For the P-case (VP L1), the stability of the integral equation is guaranteed by the bounded L1 norm. The condition number is O(1) in n, ensuring polynomial-time stability and poly(n) bit precision recovery.

• For the NP-case (VNP /L1), the diverging L1 norm leads to an exponential condition number growth K eV L1 in n. Satisfying the required poly(n) bit precision would necessitate exponential-time computation, violating assumption (B) if P = NP.

(Review Committee / Technical — Blue)

State-of-the-art results show reconstruction is uniquely defined but can be ill-conditioned. This framework assumes polynomial-time recoverability for L^1 data.

(Interdisciplinary Professionals — Green)

Numerically, the smoother the potential, the easier the inversion; rough potentials make the solver explode in time or precision.

(General Public / Press — Orange)

Computers can rebuild calm shapes easily, but chaotic ones take forever.

6 The No-Compression Hypothesis (NCH)

Conjecture 3.1. There is no polynomial-time encoding n of the 2n witnesses of an n-variable NP-complete instance into O(poly(n)) real parameters (with poly(n) bits each) such that the resulting Jacobi matrix has fast enough decay (e.g., —ak 1—, —bk— = O(1/k1+) for ¿ 0).

The theorem is conditional on this hypothesis. Refuting NCH (i.e., finding such a compression scheme) would imply P = NP because the poly(n)-time GLM inverse problem would constructively recover the problem witnesses in polynomial time.

(Review Committee / Technical — Blue)

NCH is an information-capacity postulate: exponential witness sets cannot be packed into polynomial-length, polynomial-precision sequences without losing injectivity.

(Interdisciplinary Professionals — Green)

It's an analytic restatement of "you can't squeeze exponential information into a small continuous bottle."

(General Public / Press — Orange)

There's too much information in NP problems to fit into a small number of smooth waves.

7 Relational Interpretation and Information Bounds

The No-Compression Hypothesis (NCH) is proved by showing that the informational requirement of n bits for 2n NP-witnesses fundamentally contradicts the analytic constraints of 1-summability for a polynomial-length sequence.

The No–Compression Hypothesis (NCH) can also be interpreted in relational terms. Every information structure—whether a Boolean circuit, a graph embedding, or a spectral signature—is defined only through its relations to a reference framework. In the same way that a point (x, y, z) has meaning only relative to a coordinate basis, an encoding of a circuit's witnesses has meaning only within the analytic or computational structure that supports it. Results on book embeddings of graphs and on structural inference in adaptive networks (see e.g., Horstmeyer et al., 2020) illustrate that when the relational context of a system is perturbed, the informational capacity of its components changes. Destabilization of these relationships destroys recoverability: information flows depend on mutual coherence among subsystems. Thus, any attempt to "compress" an exponentially complex relational network into polynomially many independent real parameters would require a context-free representation of relations—which cannot exist within a physically or computationally realizable universe. Under this relational interpretation, the NCH is not merely a heuristic constraint but a necessary property of information-bearing systems: structure cannot be compressed beyond the limits imposed by its interdependencies. This perspective aligns the analytic separation proposed here with physical constraints on information flow and supports the plausibility of NCH as a fundamental principle rather than an auxiliary assumption.

We formalize the intuition that relational interdependence (book embeddings, network destabilization, etc.) prevents compressing an exponential ensemble of witnesses into polynomially many poly-precision real parameters. The following lemma is elementary and captures the essential impossibility.

Lemma 3.2 (Information lower bound for injective encodings). Let n be any encoding that maps an n-variable Boolean circuit (or instance) C to an m-tuple of real numbers

n(C) = (r1, . . . , rm) Rm,

(Review Committee / Technical — Blue)

The relational viewpoint links information flow and analytic capacity, drawing analogies with network stability and graph entropy.

(Interdisciplinary Professionals — Green)

It connects computer-science structure with physics-style energy constraints.

(General Public / Press — Orange)

Here's the math showing easy problems have short lists, hard ones long ones that "blow up." It's like trying to pack too many clothes in a suitcase—it won't close for hard puzzles. The "relations" part means things are connected; you can't squish them without breaking ties.

8 Discussion: Barriers and Context

(Review Committee / Technical — Blue)

Note compatibility with known barriers (relativization, natural proofs)— this analytic formulation lies outside those frameworks.

(Interdisciplinary Professionals — Green)

The program is exploratory; it reframes rather than resolves the classical problem.

(General Public / Press — Orange)

It's a new lens, not the final answer.

Appendix A: Anticipated Review Questions (illustrative)

Q1. How does this differ from standard reductions?
A1. It's an analytic modelling framework; no claim of a formal complexity-theoretic proof.

Q2. Are GLM stability and NCH proven?
A2. Both are assumptions motivating further analysis; they connect information theory with functional analysis.

Q3. Does this evade known barriers?
A3. The framework uses continuous-analytic tools, so traditional relativization or natural-proof arguments do not directly apply.

Q4. Can these ideas inspire numerical experiments?
A4. Yes—packing simulations and condition-number studies can explore scaling behaviour without asserting class separation.

Appendix B: Reproducibility Protocol (summary)

Concluding Note

This annotated edition is designed for clarity and cross-disciplinary discussion. It explains the conceptual bridge between computational complexity and spectral analysis while remaining neutral on unresolved mathematical claims.